

Dynamic Vehicle Routing with Stochastic Requests

Alexandre Florio

November 2, 2021



Solving Large-Scale Dynamic Vehicle Routing with Stochastic Requests in Real-Time

Joint work with Jian Zhang, Kelin Luo and Tom Van Woensel

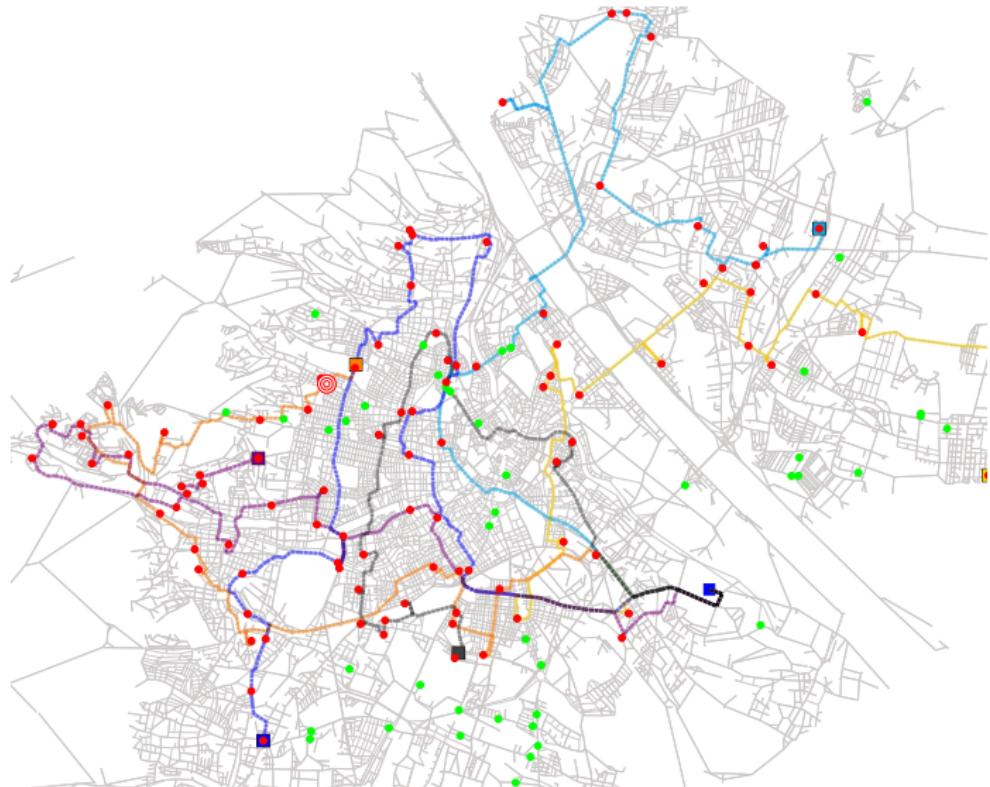


General Setting & Motivation

- Dynamic VRP with stochastic requests (DVRPSR)
 - Static (planned) and dynamic (on-demand) requests
 - Applications: technician routing, package collection (first-mile, courier or marketplace sellers)
- **Request:** pickup (or service) location and service time
 - Requests may be rejected
- Multiple drivers/vehicles
- Vehicles must return to the depot by the given deadline
- **Goal:** serve as many requests as possible
- **Contribution:** large-scale, real-time, street networks



General Setting & Motivation





Related Literature

Literature	Problem setting				Initial plan		Online policy	
	RT	Req.	Veh.	Graph	Method	Ant.	Method	Ant.
Bent & Van Hentenryck (2004)	✗	100	n/a	Synth.	MSA	✓	MSA	✓
Chen & Xu (2006)	✗	100	n/a	Synth.	CGBH	✗	CGBH	✗
Hvattum et al. (2006)	✗	130	n/a	Synth.	DSHH	✓	DSHH	✓
Ichoua et al. (2006)	✓	240	6	Synth.	TS	✓	TS	✓
Thomas (2007)	✗	50	1	Synth.	GRASP	✗	Heuristics	✓
Azi et al. (2012)	✓	144	5	Synth.	n/a		ALNS	✓
Ferrucci & Bock (2015, 2016)	✗	150	12	SSSN	TS	✓	TS	✓
Klapp et al. (2018a)	✗	40	1	Synth.	B&C	✓	Rollout	✓
Ulmer et al. (2018a)	✗	100	1	Synth.	CI	✗	VFA	✓
Ulmer et al. (2018b)	✗	100	1	Synth.	CI	✗	VFA	✓
Ulmer et al. (2019)	✗	100	1	Synth.	CI	✗	VFA+rollout	✓
van Heeswijk et al. (2019)	✗	400	n/a	Synth., SSSN	CW	✗	VFA, rollout	✓
Voccia et al. (2019)	✗	192	13	Synth.	n/a		MSA	✓
Ulmer (2020)	✓	180	3	Synth.	n/a		VFA	✓
Ulmer & Thomas (2020)	✓	50	1	Synth.	n/a		VFA	✓
This paper	✓	947	20	LSSN	PbCGBH	✓	PbPs	✓



Methodology

- Problem formulation: sequential stochastic optimization model
 - Decisions must be taken each time a request arrives
 - Note: alternative approach, “batch and optimize”
- Key contribution: approximation of the reward-to-go (potential) by relaxed knapsack models
- First decision: initial plan: column generation-based petal heuristic
- Remaining decisions: potential-based scheduling policy



Notation

$\mathcal{G} = (\mathcal{V}, \mathcal{A})$	Street network graph
\mathcal{V}	Road junctions or intersections
\mathcal{A}	Road segments
$0 \in \mathcal{V}$	Depot
d_{ij}	Length of segment $(i, j) \in \mathcal{A}$
$t_{ij} = d_{ij}/\bar{s}$	Travel time along segment $(i, j) \in \mathcal{A}$
$t(i, j)$	Duration of the fastest path from node i to node j
$[0, U]$	Service period
K	Fleet size



Notation

\mathcal{Z}	Set of static (or planned) requests
$\mathcal{D}(u)$	(Ordered) set of dynamic requests up to instant $u \in [0, U]$
(u, i, d)	Request, where $u \in [0, U]$ is the arrival/release time, $i \in \mathcal{V}$ is the location and $d \in \mathbb{R}_{>0}$ is the service duration
$T \equiv \mathcal{D}(U) $	R.v. indicating the total number of dynamic requests

We assume the spatiotemporal probability distribution of dynamic requests can be sampled from (e.g., data is available)

$\omega = \{r_1^\omega, \dots, r_{T_\omega}^\omega\}$ Sample path of the request arrival process



MDP Model: Decision Epochs and States

- Decision epoch: moment when a decision must be made
- In our setting, we have:
 - Initial decision: setup a route plan to serve all requests in \mathcal{Z}
 - Following decisions: each time a dynamic request arrives, decide whether to accept or reject the request and, if accept, how to serve it
- Hence, $1 + T$ decision epochs, where T is unknown



MDP Model: Decision Epochs and States

Definition (Route)

Sequence of pairs $\theta = ((v_0, \delta_0), \dots, (v_f, \delta_f))$, where $v_0 = v_f = 0$, $(v_i, v_{i+1}) \in \mathcal{A}$, $i \in \{0, \dots, f-1\}$, $\delta_0 = \delta_f = \emptyset$, and δ_i , $i \in \{1, \dots, f-1\}$, is the (possibly empty) set of requests to serve when arriving at the i -th node in θ

Definition (Budget of a route θ that started at instant τ)

$$b(\tau, \theta) = U - \tau - \sum_{i=0}^{f-1} t_{v_i, v_{i+1}} - \sum_{(u, i, d) \in \Delta(\theta)} d$$

where $\Delta(\theta) = \cup_{i=1}^{f-1} \delta_i$



MDP Model: Decision Epochs and States

Definition (Vehicle state)

State of vehicle $k \in \{1, \dots, K\}$ is given by

$$V_k = \begin{cases} \emptyset & \text{if } k \text{ is idle (stationed at the depot)} \\ (\tau_k, \theta_k) & \text{if } k \text{ started to travel along route } \theta_k \text{ at instant } \tau_k \end{cases}$$

Definition (State)

State S_0 (initial state) represents all parameters of the problem

State S_t , $t \in \{1, \dots, T\}$, is a $(K + 1)$ -tuple $S_t = (V_1, \dots, V_K, r_t)$, where V_k are vehicle states, and $r_t = (u_t, i_t, d_t)$ is the t -th element of $\mathcal{D}(U)$



Offline Decision: Initial Route Plan

- The decision at state S_0 is a set of routes $\{\theta_1, \dots, \theta_{k'}\}$, $k' \leq K$, such that

$$b(0, \theta_1), \dots, b(0, \theta_{k'}) \geq 0$$

$\Delta(\theta_1), \dots, \Delta(\theta_{k'})$ is a partition of \mathcal{Z}

- We let \mathcal{F} represent the set of feasible decisions at state S_0
 - Same feasible space as in a duration-constrained VRP
- Ideally, decision at S_0 should leave the system in a good condition to accept dynamic requests

At each state S_t , $t \in \{1, \dots, T\}$, a scheduling policy prescribes:

- 1 **Acceptance** decision: accept or reject request r_t
- 2 **Assignment** decision: if r_t is accepted, assign r_t to a vehicle k
- 3 **Routing** decision: when r_t is assigned to vehicle k , define how route θ_k is adjusted (in case $V_k = (\tau_k, \theta_k)$) or initialized (in case $V_k = \emptyset$) to accommodate r_t

Note: reassignment of requests among vehicles is not allowed

Consider a request $r = (u, i, d)$ and a vehicle state $V = (\tau, \theta)$

Definition (Cheapest Insertion (CI) Routing Policy ρ_{CI})

Under the CI routing policy ρ_{CI} , route θ is adjusted into a new route $\rho_{\text{CI}}(\theta, r)$ in such a way that $\Delta_u(\rho_{\text{CI}}(\theta, r)) \setminus \Delta_u(\theta) = \{r\}$, $\Delta_u(\theta)$ is a subsequence of $\Delta_u(\rho_{\text{CI}}(\theta, r))$ and the budget $b(\tau, \rho_{\text{CI}}(\theta, r))$ is maximized

→ Polynomial time

Definition (Reoptimization Routing Policy ρ_{R})

Under the reoptimization routing policy ρ_{R} , route θ is adjusted into a new route $\rho_{\text{R}}(\theta, r)$ in such a way that $\Delta_u(\rho_{\text{R}}(\theta, r)) \setminus \Delta_u(\theta) = \{r\}$ and the budget $b(\tau, \rho_{\text{R}}(\theta, r))$ is maximized

→ NP-hard



Online Decisions: State Transitions

- Let \mathcal{X}_t be the set of possible decisions when at state S_t
- A policy π maps each possible state S_t to a decision $X^\pi(S_t) \in \mathcal{X}_t$
- State transition:

$$S_{t+1} = S^M(S_t, x_t, r_{t+1})$$

- Decisions must be taken in real-time (i.e., in seconds, not minutes)



MDP Model: Rewards and Objective

Rewards:

$$R(S_t, x) = \begin{cases} 1 & \text{if } x \in \mathcal{X}_t \text{ is an 'accept' decision} \\ 0 & \text{otherwise} \end{cases}$$

Objective:

$$\max_{\pi \in \Pi, y \in \mathcal{F}} \mathbb{E} \left[\mathbb{E} \left[\sum_{t=1}^T R(S_t, X^\pi(S_t)) \middle| S_1 \right] \right]$$

where Π is the set of feasible policies and \mathcal{F} is the set of feasible initial plans

Online decisions:

- **Potential-based policy (PbP)**
- **Simplified PbP (S-PbP)**
- Benchmark policies:
 - Greedy policy with cheapest insertion (CI) and reoptimization (GP_{CI} and GP_R)
 - PFA policies with CI and reoptimization (PFA_{CI} and PFA_R)
 - Rollout policies with GP_{CI} and PFA_{CI} as base policies
(R_R - $GP_{CI}(H)$ and R_R - $PFA_{CI}(H)$)

Offline decisions:

- Myopic plan
- **Potential-based plan**



Potential of a State

Given a policy π , the potential of a state S_t is the expected reward-to-go:

$$\Phi_{\pi}(S_t) = \mathbb{E} \left[\sum_{t'=t}^T R(S_{t'}, X^{\pi}(S_{t'})) \middle| S_t \right] \quad t \neq 0$$



At state S_t , let $x_t^- \in \mathcal{X}_t$ be the decision that rejects request r_t

Decision Rule 1 (Acceptance Rule)

The optimal policy $\pi^* \in \Pi$ accepts request r_t if and only if there exists an 'accept' decision $x_t^+ \in \mathcal{X}_t \setminus \{x_t^-\}$ such that:

$$1 + \mathbb{E} \left[\Phi_{\pi}(S^M(S_t, x_t^+, r_{t+1})) \right] \geq \mathbb{E} \left[\Phi_{\pi}(S^M(S_t, x_t^-, r_{t+1})) \right]$$



Decision Rule 2 (Assignment Rule)

Let $\{\mathcal{X}_t^k\}$ be a partition of $\mathcal{X}_t \setminus \{x_t^-\}$, in which \mathcal{X}_t^k is the set of decisions where r_t is accepted and assigned to vehicle k . Also, for all $\mathcal{X}_t^k \neq \emptyset$, let

$$x_t^k = \arg \max_{x_t \in \mathcal{X}_t^k} \mathbb{E} \left[\Phi_\pi(S^M(S_t, x_t, r_{t+1})) \right]$$

Then, the optimal policy π^* assigns an accepted request r_t to vehicle

$$k^* = \arg \max_k \mathbb{E} \left[\Phi_\pi(S^M(S_t, x_t^k, r_{t+1})) \right]$$

Idea: approximate potentials and prescribe decisions according to DR1 and DR2



Towards an Approximation Model: Two Remarks

Remark 1 (Late Depot Arrival)

Since the goal is to maximize the number of accepted requests, drivers return to the depot at an instant near the end of the service period

Remark 2 (Request Order Preservation)

The relative order of scheduled requests does not change (CI policy), or changes only slightly (reopt. policy), when a new request is assigned to a route



Modeling Insight: Effective Vehicle Speed

Remarks 1 and 2 suggest the concept of **effective speed**:

Definition (Effective Vehicle Speed)

The effective speed of a vehicle k with state $V_k = (\tau_k, \theta_k)$ is the speed such that k arrives at the depot exactly at instant U :

$$\check{s}_{u_t}(V_k) = \frac{d_{u_t}(V_k)\bar{s}}{d_{u_t}(V_k) + b(\tau_k, \theta_k)\bar{s}}$$

where u_t is the current time instant and $d_{u_t}(V_k)$ is the remaining distance to be traveled along route θ_k at instant u_t

The effective speed allows us to predict the location of a vehicle at any future instant



Modeling Insight: Assignment Costs

$\mathcal{V}_{u_t}(V_k, u')$ (Predicted) set of nodes along θ_k not yet traversed by vehicle $V_k = (\tau_k, \theta_k)$ at a future instant u' , obtained by simulating route θ_k under the effective speed

The minimum cost (or budget consumption) when assigning a future request $r' = (u', i', d')$ to vehicle k is therefore:

$$c_{u_t}(V_k, r') = \min_{j \in \mathcal{V}_{u_t}(V_k, u')} t(j, i') + t(i', j) + d'$$



Multiple-Knapsack Approximation of the Potential

Given H sample paths $\Omega = \{\omega_1, \dots, \omega_H\}$ for the remaining horizon:

$$\mathbb{E}[\Phi_{\pi^*}(S_{t+1})|x_t] \approx \hat{\Phi}_{\pi^*}(S_{t+1}|x_t) = \frac{1}{H} \sum_{\omega \in \Omega} \phi_{\pi^*}^\omega(S_{t+1}|x_t)$$

where

$$\begin{aligned}\phi_{\pi^*}^\omega(S_{t+1}|x_t) &= \max \quad \sum_{k \in \bar{K}} \sum_{r \in \omega} z_{kr} \\ \text{s.t.} \quad &\sum_{r \in \omega} c_{u_t}(V_k, r) z_{kr} \leq b(\tau_k, \theta_k) \quad k \in \bar{K} \\ &\sum_{k \in \bar{K}} z_{kr} \leq 1 \quad r \in \omega \\ &0 \leq z_{kr} \leq 1 \quad k \in \bar{K}, r \in \omega\end{aligned}$$



Potential-based Policy (PbP, short version)

Given routing policy $\rho \in \{\rho_{\text{Cl}}, \rho_{\text{R}}\}$. Upon arrival of $r_t = (u_t, i_t, d_t)$:

- 1 If there is an idle vehicle k that can serve request r_t , accept r_t and assign it to vehicle k (**same step for all policies**)
- 2 Otherwise, initialize a set of candidate decisions $\tilde{\mathcal{X}}_t = \{\tilde{x}_t^-\}$ with the 'reject' decision \tilde{x}_t^- related to request r_t
- 3 If it is feasible to serve r_t with vehicle k under a given routing policy ρ , add the corresponding decision \tilde{x}_t^k to set $\tilde{\mathcal{X}}_t$
- 4 Compute $\hat{\Phi}_{\pi^*}(S_{t+1}|\tilde{x}_t)$ for each $\tilde{x}_t \in \tilde{\mathcal{X}}_t$, and select the best decision according to Decision Rules 1 and 2



Potential-based Policy (PbP, long version)

Given routing policy $\rho \in \{\rho_{\text{CI}}, \rho_{\text{R}}\}$. Upon arrival of $r_t = (u_t, i_t, d_t)$:

- 1 If there is an idle vehicle k that can serve request r_t , r_t is accepted and assigned to vehicle k ; route θ_k is initialized to the fastest route from the depot to node i_t , and back (**same step for all policies**)
- 2 Otherwise, a set $\tilde{\mathcal{X}}_t = \{\tilde{x}_t^-\}$ of candidate decisions is initialized, where \tilde{x}_t^- is the 'reject' decision related to request r_t
- 3 For each non-idle vehicle $V_k = (\tau_k, \theta_k)$, we compute the adjusted route $\rho(\theta_k, r_t)$. If $b(\tau_k, \rho(\theta_k, r_t)) \geq 0$, then the decision \tilde{x}_t^k of accepting and assigning r_t to k under routing policy ρ is feasible and is added to $\tilde{\mathcal{X}}_t$
- 4 A decision is selected from $\tilde{\mathcal{X}}_t$ by evaluating Decision Rules 1 and 2 based on approximations $\hat{\Phi}_{\pi^*}(S_{t+1}|\tilde{x}_t)$ of $\mathbb{E}[\Phi_{\pi^*}(S_{t+1})|\tilde{x}_t]$, $\tilde{x}_t \in \tilde{\mathcal{X}}_t$, computed by the multiple-knapsack model



Simplified Potential-based Policy (S-PbP)

- Policy PbP requires solving the multiple-knapsack model up to $H(K + 1)$ times
 - Depending on our real-time requirements, this may take too long
 - Note: algorithm can be easily parallelized (not implemented)
- S-PbP: trades-off approximation accuracy for computational efficiency



Simplified Potential-based Policy (S-PbP)

Consider the following single-knapsack model:

$$\begin{aligned} p_{u_t}^\omega(V_k) = \max \quad & \sum_{r \in \omega} z_r \\ \text{s.t.} \quad & \sum_{r \in \omega} c_{u_t}(V_k, r) z_r \leq b(\tau_k, \theta_k) \\ & 0 \leq z_r \leq 1 \quad r \in \omega \end{aligned}$$

Problem: considers vehicles in isolation; does not take into account competition of vehicles for requests, so the potential is overestimated



Simplified Potential-based Policy (S-PbP)

To compensate for the overestimation bias, we define the ratio α^ω :

$$\alpha^\omega = \frac{\phi_{\pi^*}^\omega(S_{t+1} | \tilde{x}_t^-)}{\sum_{V_k \neq \emptyset} p_{u_t}^\omega(V_k)}$$

Policy S-PbP approximates the potential as follows:

$$\mathbb{E}[\Phi_{\pi^*}(S_{t+1}) | \tilde{x}_t^k] \approx \frac{1}{H} \sum_{\omega \in \Omega} \alpha^\omega \left(p_{u_t}^\omega(V_k^+) + \sum_{V_l \neq \emptyset, l \neq k} p_{u_t}^\omega(V_l) \right)$$

where V_k^+ is the state of vehicle k after 'accept' decision \tilde{x}_t^k

Complexity: requires solving the multiple-knapsack model H times, and solving the single-knapsack model up to $2HK$ times



Benchmark Policies: Greedy Policy

Given routing policy $\rho \in \{\rho_{\text{CI}}, \rho_{\text{R}}\}$. Upon arrival of $r_t = (u_t, i_t, d_t)$:

- 1 If there is an idle vehicle k that can serve request r_t , accept r_t and assign it to vehicle k (**same step for all policies**)
- 2 Otherwise, request r_t is accepted whenever there is a non-idle vehicle k , $V_k = (\tau_k, \theta_k)$, such that $b(\tau_k, \rho(\theta_k, r_t)) \geq 0$
- 3 If r_t is accepted, assign r_t to the non-idle vehicle k , $V_k = (\tau_k, \theta_k)$, such that $b(\tau_k, \theta_k) - b(\tau_k, \rho(\theta_k, r_t))$ is minimum

PFA-based policy approximates the potential as follows:

$$\hat{\Phi}_{\text{PFA}}(S_{t+1}|x_t) = \begin{cases} \gamma \sum_{V_k \neq \emptyset} b(\tau_k, \theta_k) & \text{if } x_t = x_t^- \in \mathcal{X}_t \\ \gamma \left(b(u_t, \rho(\theta_k, r_t)) + \sum_{i \neq k, V_i \neq \emptyset} b(\tau_i, \theta_i) \right) & \text{if } x_t \in \mathcal{X}_t^k \end{cases}$$

Parameter γ :

- Interpretation: expected reward per unit of budget
- Trained offline by grid search and simulation



Benchmark Policies: Rollout Policies

Upon arrival of request $r_t = (u_t, i_t, d_t)$:

- 1 If there is an idle vehicle k that can serve request r_t , accept r_t and assign it to vehicle k (**same step for all policies**)
- 2 Otherwise, initialize a set of candidate decisions $\tilde{\mathcal{X}}_t = \{\tilde{x}_t^-\}$ with the 'reject' decision \tilde{x}_t^- related to request r_t
- 3 If it is feasible to serve r_t with vehicle k under a given routing policy ρ , add the corresponding decision \tilde{x}_t^k to set $\tilde{\mathcal{X}}_t$
- 4 Simulate each candidate decision over H sample paths under a given base policy; then, prescribe the decision with the highest simulated reward



Offline Decision: Initial Route Plan

Set-partitioning formulation:

$$\begin{aligned} & \max_{\{\theta_1, \dots, \theta_{k'}\} \subset \Theta, k' \leq K} g(S_1) \equiv \mathbb{E} [\Phi_{\pi^*}(S_1)] \\ \text{s.t. } & \sum_{k \in \{1, \dots, k'\}} [(u, i, d) \in \Delta(\theta_k)] = 1 \quad (u, i, d) \in \mathcal{Z} \\ & V_k = (0, \theta_k) \quad k \in \{1, \dots, k'\} \\ & V_k = \emptyset \quad k \in \{k' + 1, \dots, K\} \end{aligned}$$

where $S_1 = (V_1, \dots, V_K, r_1)$, and $\Theta = \{\theta : b(0, \theta) \geq 0\}$ is the set of feasible routes



Offline Decision: Initial Route Plan

Myopic plan: maximize budget

$$\tilde{g}_{\text{myo}}(S_1) = \sum_{V_k \neq \emptyset} b(0, \theta_k) + \sum_{V_k = \emptyset} U$$

- Rational: more budget means more time to serve dynamic requests
- Problem reduces to a duration-constrained VRP
- Column generation-based petal heuristic



Offline Decision: Initial Route Plan

Potential-based plan: maximize potential

$$\tilde{g}_{\text{pb}}(S_1) = \frac{1}{H} \sum_{\omega \in \Omega} \sum_{V_k \neq \emptyset} p_0^\omega(V_k)$$

- Rational: serve all static requests by dispatching drivers along promising routes
- Non-linear cost function due to competition among drivers for requests; approximation via single-knapsack model
- Petal heuristic: reevaluate costs of routes generated by the DCVRP algorithm and resolve set-partitioning model



Computational Study

- Real street network of Vienna (16,080 nodes and 36,424 arcs)
- Service period: 10 hours
- Number of vehicles: $K \in \{2, 3, 5, 6, 10, 12, 20\}$
- Request rate (per minute): $\Lambda \in \{0.2, 0.4, 0.8, 1.5\}$
 - Smallest instances larger than most instances from previous works
- Degree of dynamism: $\eta \in \{75\%, 85\%, 90\%, 95\%\}$
- Three spatiotemporal request distributions:
 - Uniform, time-independent (UTI)
 - Clustered, time-independent (CTI)
 - Clustered, time-dependent (CTD)
- A modular DVRPSR simulator is implemented in C++



Real World Street Network





Instances, Policies and Offline Planners Simulated

Instance parameters					Algorithms		Simulations
Λ	η	K	Distribution(s)	Scenario(s)	Offline	Online	
0.2	0.75	2, 3	UTI, CTI, CTD	5	MY, PB	GP_{cl} , GP_R , PFA_{cl} , PFA_R , R_{cl} - $GP_{cl}(10, 25, 50, 100)$, R_R - $GP_{cl}(10, 25, 50, 100)$, R_R - $PFA_{cl}(25, 50, 100)^*$, S-PbP, PbP	12,270
0.4	0.85	3, 5	UTI, CTI, CTD	5	MY, PB	GP_{cl} , GP_R , PFA_{cl} , PFA_R , R_{cl} - $GP_{cl}(10, 25, 50, 100)$, R_R - $GP_{cl}(10, 25, 50, 100)$, R_R - $PFA_{cl}(25, 50, 100)^*$, S-PbP, PbP	12,270
0.8	0.90	6, 12	UTI, CTI, CTD	5	PB	GP_{cl} , GP_R , PFA_{cl} , PFA_R , R_{cl} - $GP_{cl}(10)$, R_R - $GP_{cl}(10)$, R_R - $PFA_{cl}(10)$, S-PbP, PbP	3,210
1.5	0.95	10, 20	UTI	1	PB	GP_{cl} , GP_R , PFA_{cl} , PFA_R , R_{cl} - $GP_{cl}(10)$, R_R - $GP_{cl}(10)$, R_R - $PFA_{cl}(10)$, S-PbP, PbP	214

Notes. MY, myopic planner; PB, potential-based planner; the online algorithms marked by * are tested with the potential-based planner only.



Potential-based vs Myopic Plans

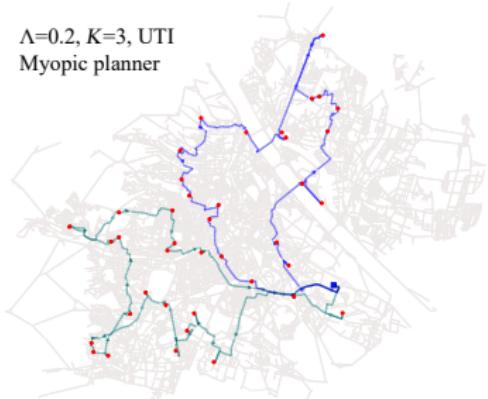
Average number of accepted requests:

Policy	Myopic Plan	Potential-based Plan	Diff
GP _{CI}	54.1	61.2	13.0%
GP _R	54.9	62.7	14.1%
R _{CI} -GP _{CI} (10)	63.0	66.2	5.1%
R _{CI} -GP _{CI} (25)	65.8	69.4	5.4%
R _{CI} -GP _{CI} (50)	67.3	71.1	5.6%
R _{CI} -GP _{CI} (100)	68.2	72.1	5.8%
R _R -GP _{CI} (10)	63.4	67.2	6.1%
R _R -GP _{CI} (25)	66.2	70.7	6.7%
R _R -GP _{CI} (50)	67.7	72.2	6.8%
R _R -GP _{CI} (100)	68.5	73.3	7.0%
S-PbP(50)	72.0	75.2	4.4%
PbP(50)	72.9	77.2	5.9%
PFA _{CI}	61.3	71.8	17.1%
PFA _R	62.1	72.6	16.9%
Avg	66.4	72.4	9.0%

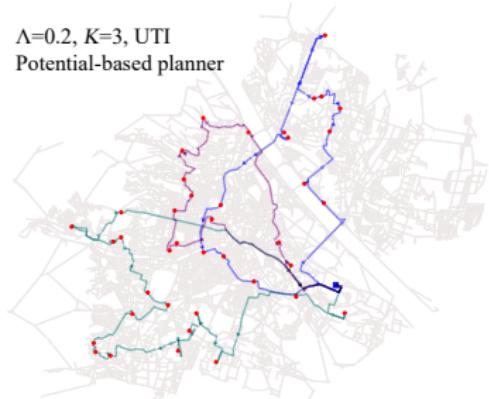


Potential-based vs Myopic Plans

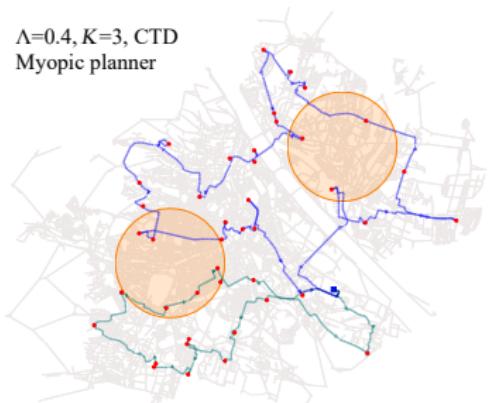
$\Lambda=0.2, K=3, \text{UTI}$
Myopic planner



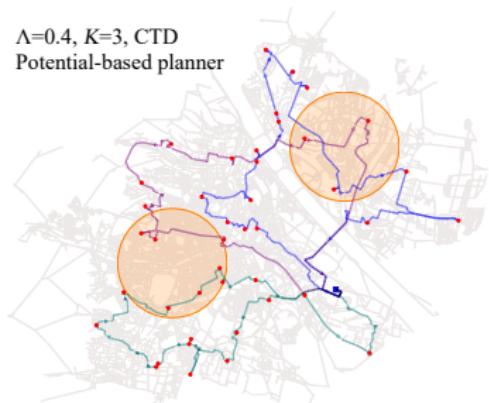
$\Lambda=0.2, K=3, \text{UTI}$
Potential-based planner



$\Lambda=0.4, K=3, \text{CTD}$
Myopic planner

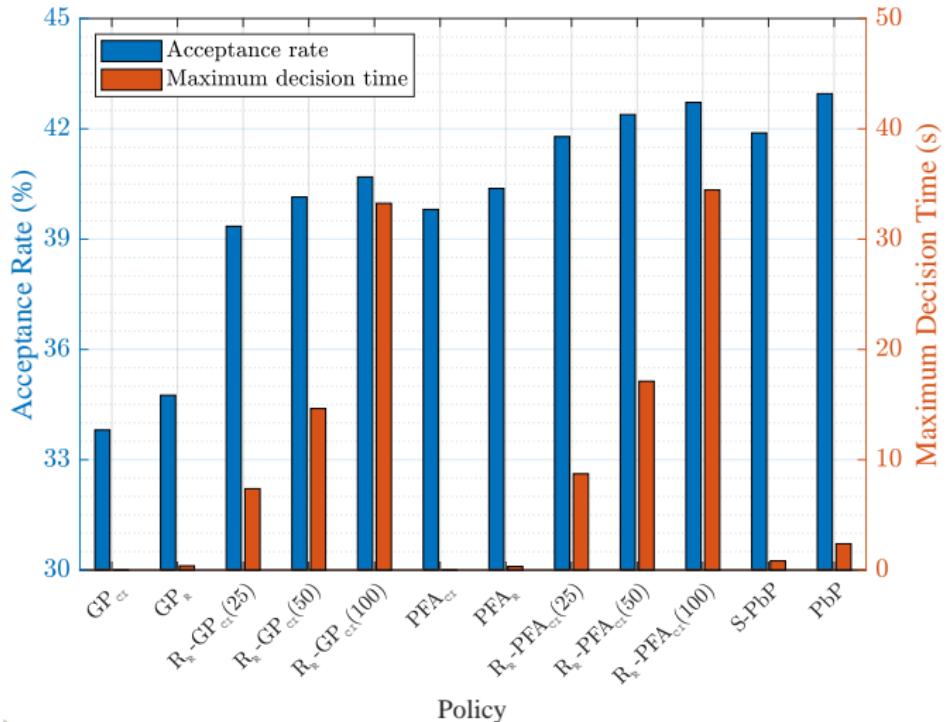


$\Lambda=0.4, K=3, \text{CTD}$
Potential-based planner



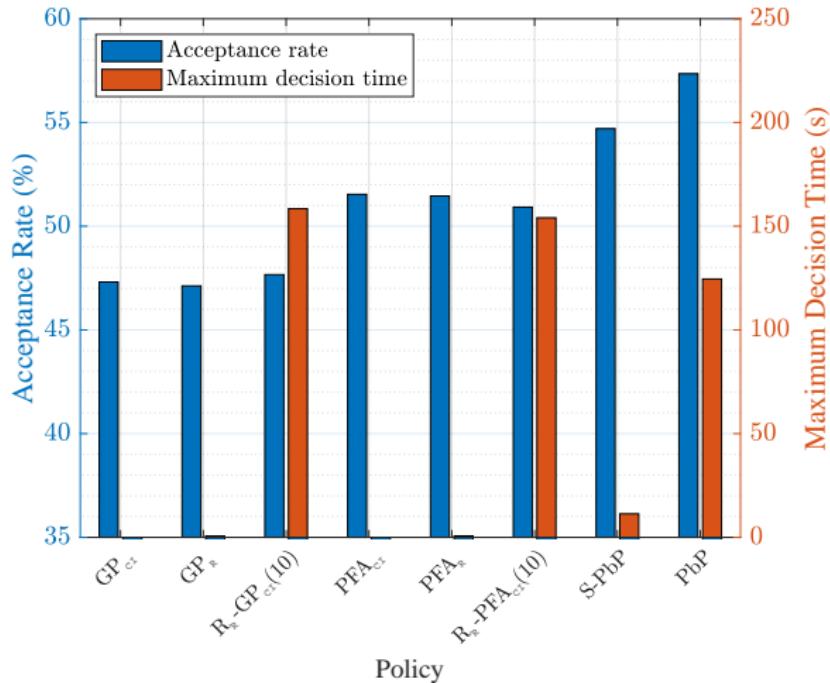


Policy Comparison ($\Lambda \in \{0.2, 0.4\}$)



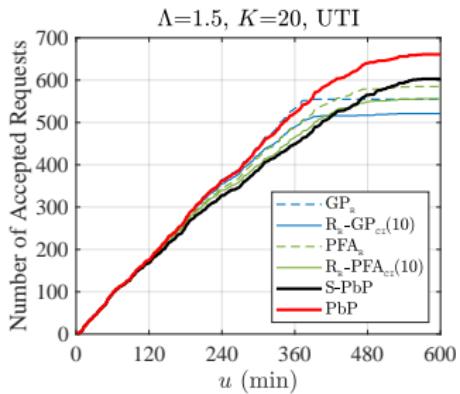
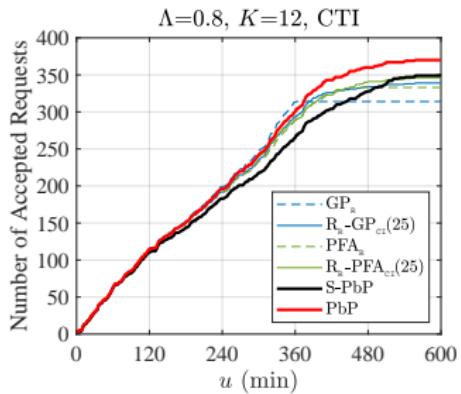
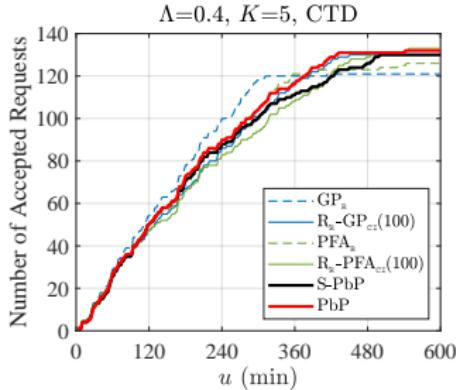
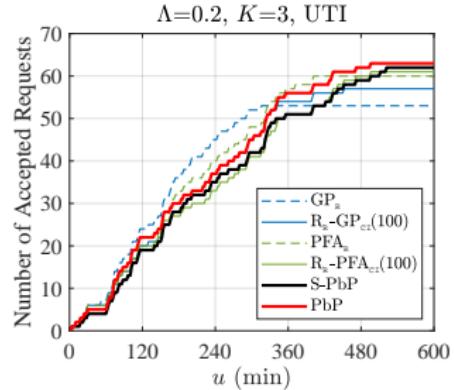


Policy Comparison ($\Lambda \in \{0.8, 1.5\}$)

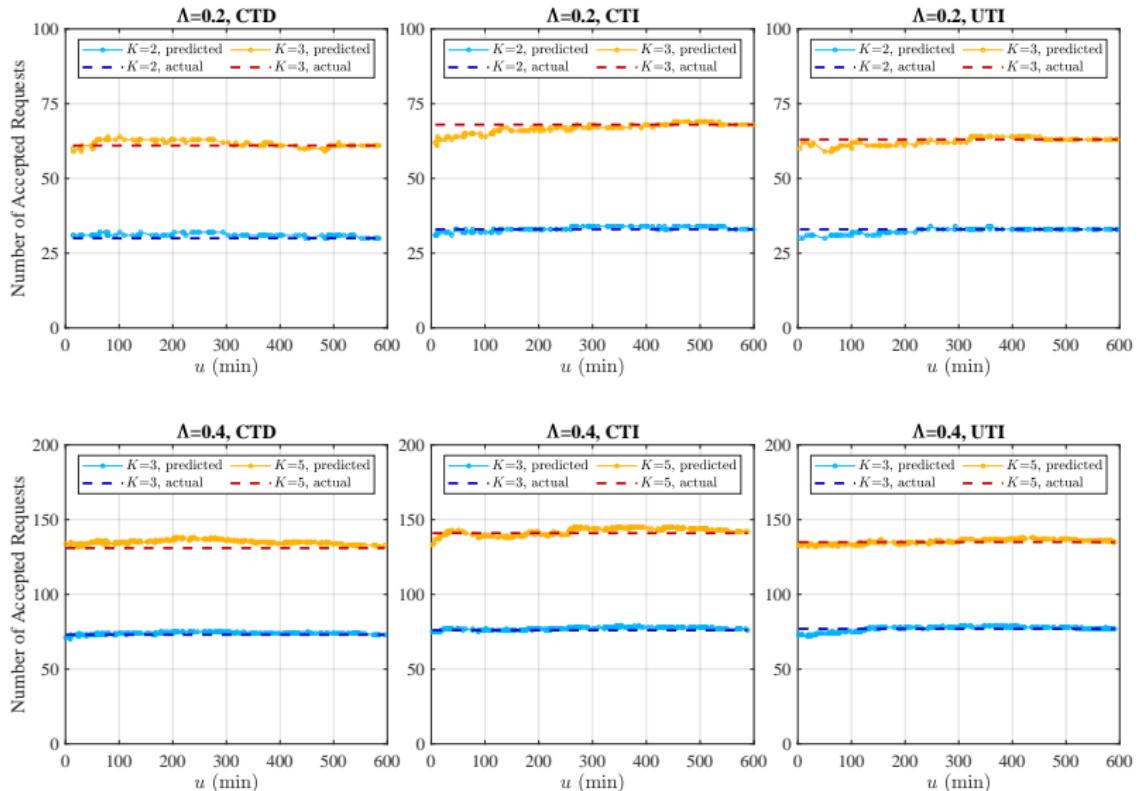




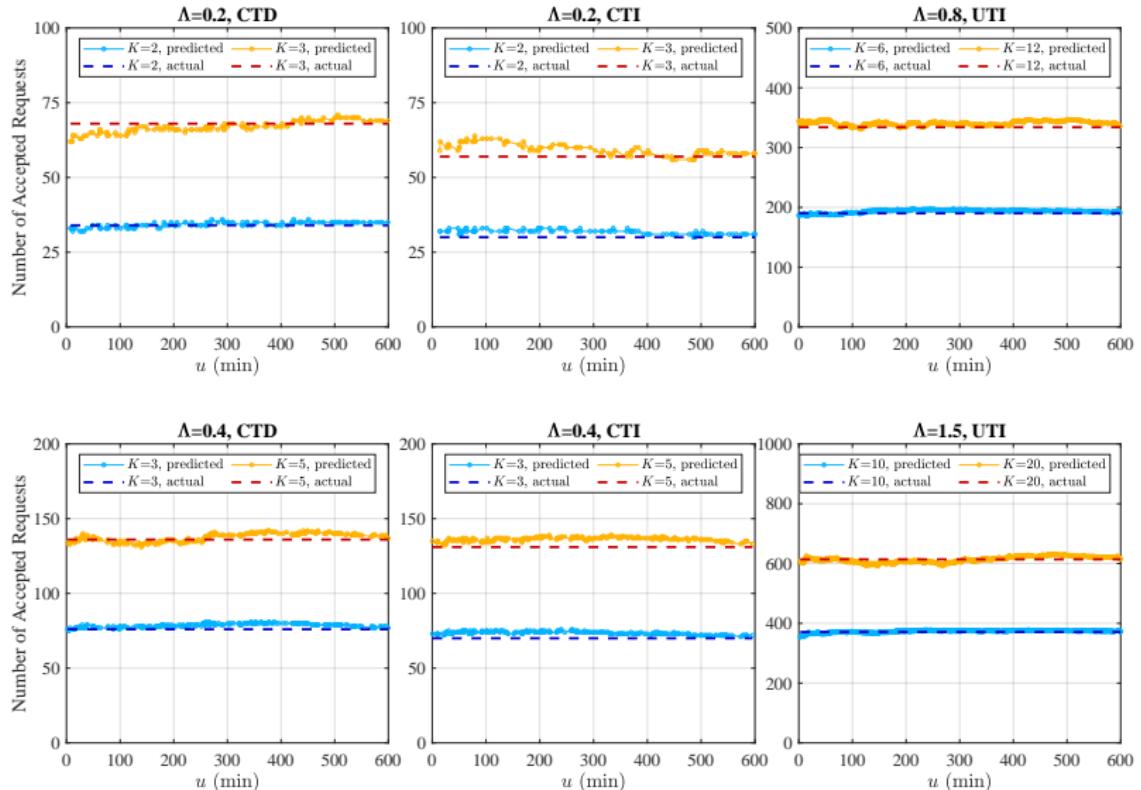
Acceptance Profiles



Multiple-Knapsack Potential Approximations



Single-Knapsack Potential Approximations





Conclusions

- Main contributions and takeaways:
 - Expected reward-to-go can be accurately and efficiently approximated by knapsack models
 - Accurate potential approximations enable high-performing scheduling policies, which outperform classical ADP methods traditionally used for DVRPs such as rollout algorithms and PFA
 - Potential is key for both offline and online decisions
 - Coverage of the service area is more important than budget alone
- Possible extensions and future research:
 - Vehicle capacity constraints
 - Time windows
 - Self-imposed time windows (i.e., time window assignment)
 - Pickup and delivery
 - Reassignment of requests among vehicles (more complex)
 - Comparison against 'batch and optimize' approach



Acknowledgements

This project has received funding from the European Union's Horizon 2020 research and innovation programme under the Marie Skłodowska-Curie grant agreement No 754462

Computational experiments were performed on the Dutch national e-infrastructure with the support of SURF Cooperative

Geographical data for Vienna are copyrighted to OpenStreetMap contributors and available at <http://openstreetmap.org>

Questions & Discussions